

Lecture 10:

CPU – exception

Hunjun Lee


hunjunlee@hanyang.ac.kr

CPUs must prepare for unexpected events

- ◆ The CPUs must detect and handle unexpected conditions
 - Instructions may fail and cannot complete
 - External I/O devices
- ◆ Option #1: The SW continuously checks every possible unexpected events (i.e., polling)
 - Acceptable for simple embedded systems
- ◆ Option #2: The HW dynamically detects unexpected events
 - Transparently transfer control to an exception handler (that knows how to resolve the condition and move back to your program)

Interrupt control transfer

- ◆ Unlike a normal function call, the interrupted thread (we'll get to this later) does not anticipate control transfer
- ◆ Control should be later returned to the main thread
- ◆ The control transfer should be 100% transparent to the interrupted thread



Types of Interrupts

◆ Exceptions: Synchronous interrupts

- Caused by (Tied to) a specific instruction
 - Illegal opcode, div by zero, arithmetic overflow, page fault ...
- **Should be handled immediately!**

◆ Interrupts: Asynchronous interrupts

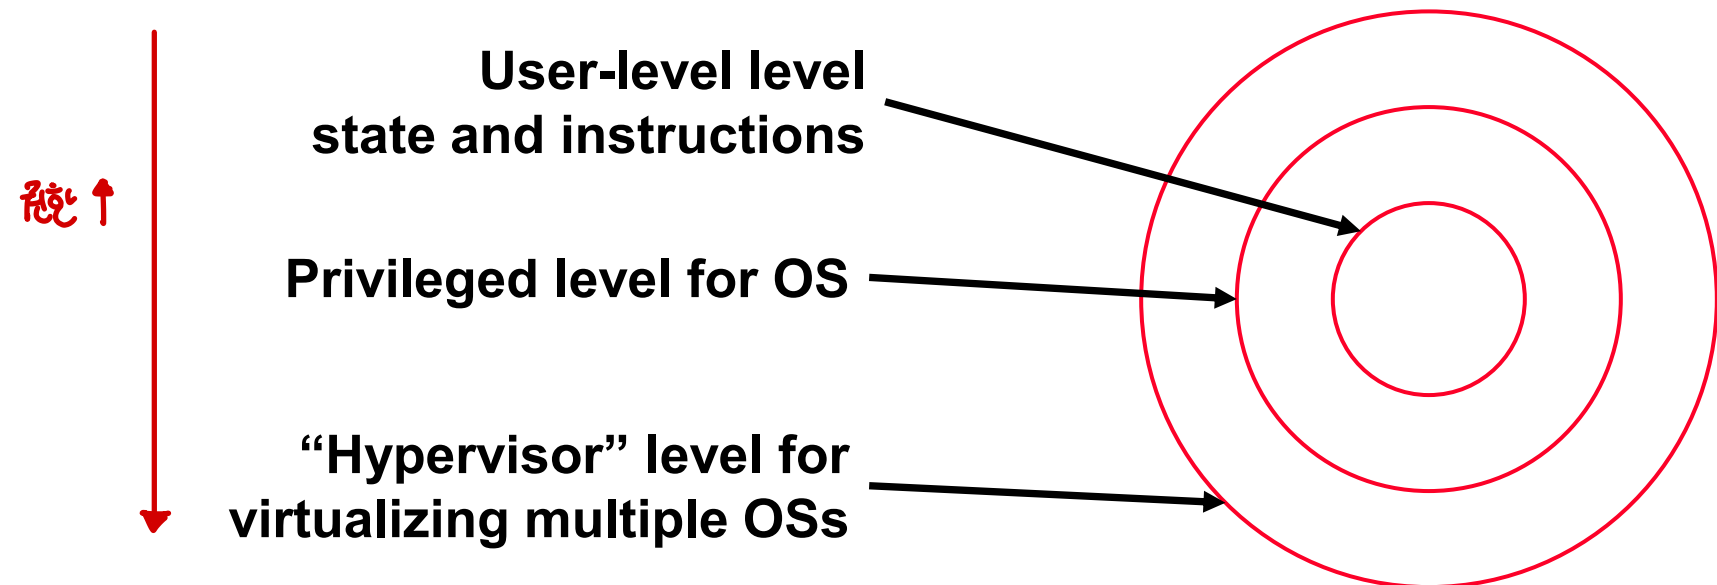
- External events (not tied to a particular instruction)
 - Ex) I/O events, power failure, ...
- **Can be postpone to some extent**
 - However, there exists a priority depending on the type

◆ System call/trap instruction

- An instruction whose only purpose is **to raise an exception** (e.g., print ...)



Privilege levels

- ◆ The OS must somehow be more powerful to ensure transparency!
- ◆ There are multiple privilege modes depending on what should be handled!




Virtualization

- ◆ Virtualization is a software to create **virtual representations of servers, storage, networks, and other physical machines**
 - This enables the software to control the number of allocated resources



Privilege transfer


- ◆ User-level code does not run in the privileged mode
- ◆ CPU enters the privileged mode on interrupts (user code transfers control to the OS kernel)
- ◆ The handler restores privilege level back to the user mode (before restarting the user code)



Precise interrupts & execution

Interrupt 발생은 키보드가 보냈을 Interrupt

- ◆ A precise interrupt appears to take place “exactly” in between two instructions
 - Older instructions complete
 - Younger instructions are flushed (as if never happened)
- ◆ For synchronous interrupts, execution stops just before the faulting instruction, and resume after handling is done



Supporting a precise exception

- ◆ When an instruction is detected to have caused an exception, there should be the following mechanisms
 - Flushes all younger instructions
 - Keep the architectural states precise (register file, PC, memory)
 - Saves PC and registers
 - Redirects the execution to the appropriate exception handling instructions

Why precise exceptions?

- ◆ Because the ISA says so ...
 - The programmer thinks that the instructions change the architectural states in order
- ◆ For debugging ...
 - You'll know why if you have ever debugged a multithread program
- ◆ Easy recovery from exceptions
 - We do not need to change the states when handling exceptions
- ◆ Easy to restart the process after the exception
- ◆ Also, you can think of syscall (e.g., print)

Supporting a precise exception

- ◆ When an instruction is detected to have caused an exception, there should be the following mechanisms

- Flushes all younger instructions
- Keep the architectural states precise (register file, PC, memory)
 - Saves PC and registers
- Redirects the execution to the appropriate exception handling instructions

Stopping and restarting a pipeline

Exception for I_3 detected @ ID stage

	t_0	t_1	t_2	t_3	t_4	t_5	t_6	t_7	t_8	<i>Handling</i>	
IF	I_0	I_1	I_2	I_3	I_4	bub	bub	I_h	I_{h+1}	I_{h+2}	I_{h+3}
ID		I_0	I_1	I_2	I_3	bub	bub	bub	I_h	I_{h+1}	I_{h+2}
EX			I_0	I_1	I_2	bub	bub	bub	bub	I_h	I_{h+1}
MEM				I_0	I_1	I_2	bub	bub	bub	bub	I_h
WB					I_0	I_1	I_2	bub	bub	bub	bub

How would things look different for asynchronous interrupts

Exception sources in different stages

◆ IF

- Instruction memory address/protection fault

◆ ID

- Illegal opcode → 하드웨어로 동작할 수 없는 명령어 소프트웨어 실행
- Trap to SW emulation of unimplemented instructions
- System call instruction (an intended exception requested by SW)

◆ EX

- Invalid results: overflow, divide-by-zero, ...

◆ MEM

- Data memory address, page fault

◆ WB

- There is no exception by now ...

Stopping and restarting a pipeline

Exception for I_3 detected @ ID stage

	t_0	t_1	t_2	t_3	t_4	t_5	t_6	t_7	t_8	<i>Handling</i>	
IF	I_0	I_1	I_2	I_3	I_4	bub	I_h	I_{h+1}	I_{h+2}	I_{h+3}	I_{h+4}
ID		I_0	I_1	I_2	I_3	bub	bub	I_h	I_{h+1}	I_{h+2}	I_{h+3}
EX			I_0	I_1	I_2	bub	bub	bub	I_h	I_{h+1}	I_{h+2}
MEM				I_0	I_1	I_2	bub	bub	bub	I_h	I_{h+1}
WB					I_0	I_1	I_2	<i>does not cause exception</i>			

Supporting a precise exception

- ◆ When an instruction is detected to have caused an exception, there should be the following mechanisms
 - Flushes all younger instructions

- Keep the architectural states precise (register file, PC, memory)
 - Saves PC and registers
- Redirects the execution to the appropriate exception handling instructions

Interrupt handling @ MIPS - 1

◆ Save some architectural states before calling an interrupt handler

↓ Handler 속 저장 필요

- Saves PC to the Exception program counter (EPC)
- Saves r26, r27 registers (\$k0~\$k1: reserved for OS kernel)

◆ How to know the cause of the interrupt

원인

- CPU records the cause of the interrupt in a privileged registers
- **Option #1)** Use cause register to indicate the problem
 - The initial handler uses the register to jump to another handler at the kernel space (flexible, but slow) ← Cause Register 값에 따라 다른 handler로
- **Option #2)** Jump to different address depending on the cause
 - The HW determines which instruction to jump to (inflexible, but fast)
 - Undefined opcode: 8000 0000
 - Overflow: 8000 0180


Interrupt handling @ MIPS - 2

- ◆ If the interrupt is can be corrected → take corrective action and return to the EPC
- ◆ Otherwise, terminate the program and report error


Interrupt handling @ MIPS - 3

◆ Details on interrupt handling registers

- Exception program counter (EPC, CR14)
- Interrupt cause register (CR13):
 - **What caused the interrupt?**

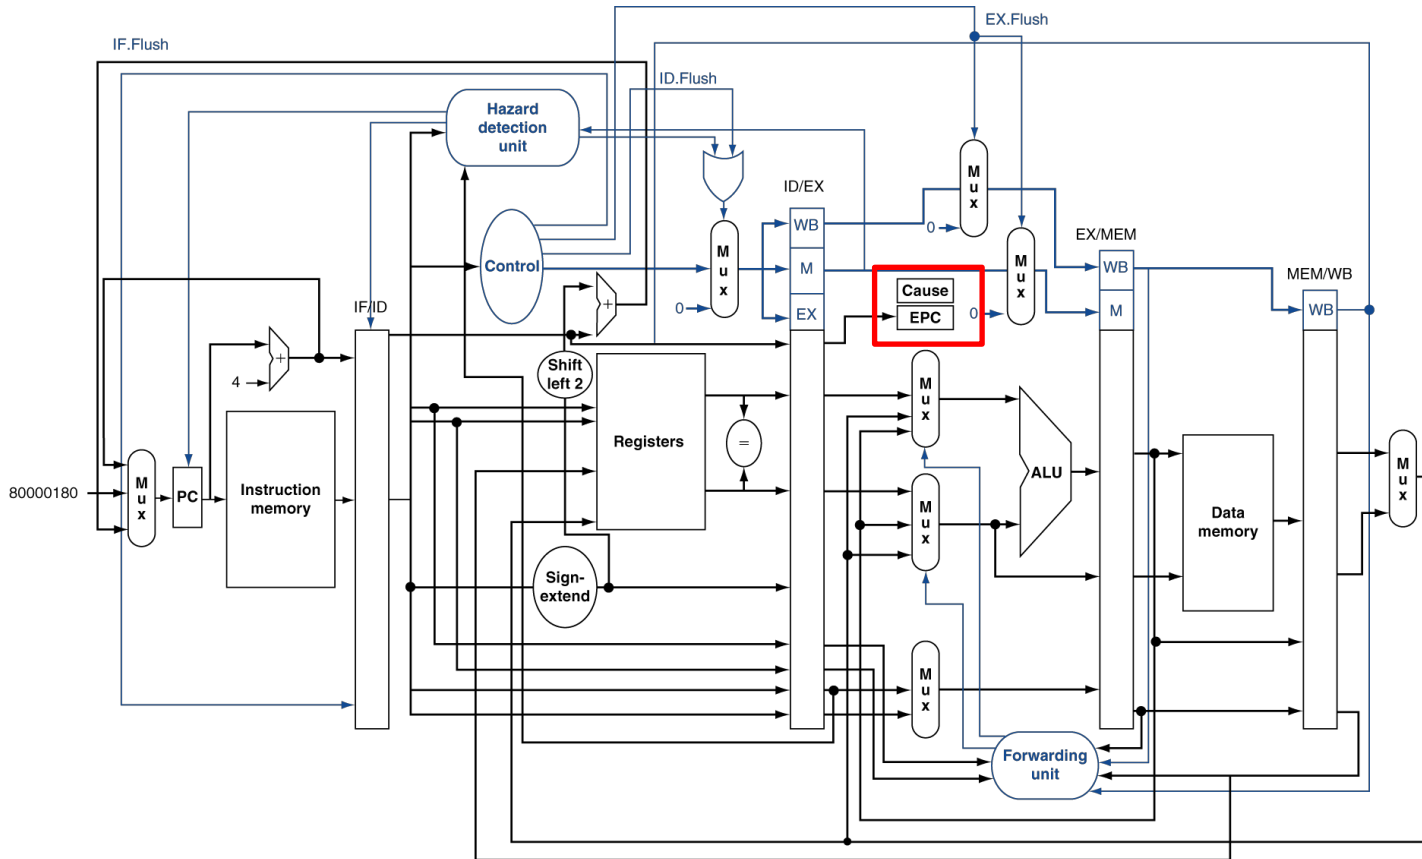


- Interrupt status register (CR12):
 - **Enable/Disable interrupts + Set privilege modes**



- Automatically set on interrupts
- Also accessed by the “move from/to register-x”
 - `mfc0 $Rx, CRx` ($Rx \leftarrow CRx$) / `mtc0 $Rx, CRx` ($CRx \leftarrow Rx$)

Pipeline with Exceptions



We need a flush mechanism till EX stage (to support exceptions at MEM stage)

→ WB에 exception X, MEM에 exception 발생하면, IF ~ EX의 instruction은 flush
→ IF ~ EX의 flush mechanism이 필요, 2차 flush는 IF ~ EX의 instruction을 flush 하도록 함.

Exception example

◆ Exception on add (overflow)

40 sub \$11, \$2, \$4

44 and \$12, \$2, \$5

48 or \$13, \$2, \$6

4C add \$1, \$2, \$1

50 slt \$15, \$6, \$7

54 lw \$16, 50(\$7)

...

→ EX0004 EXCEPTION

↓ 0/2H 27H FLUSH

Cause + EPC 1120

Jump to Handler

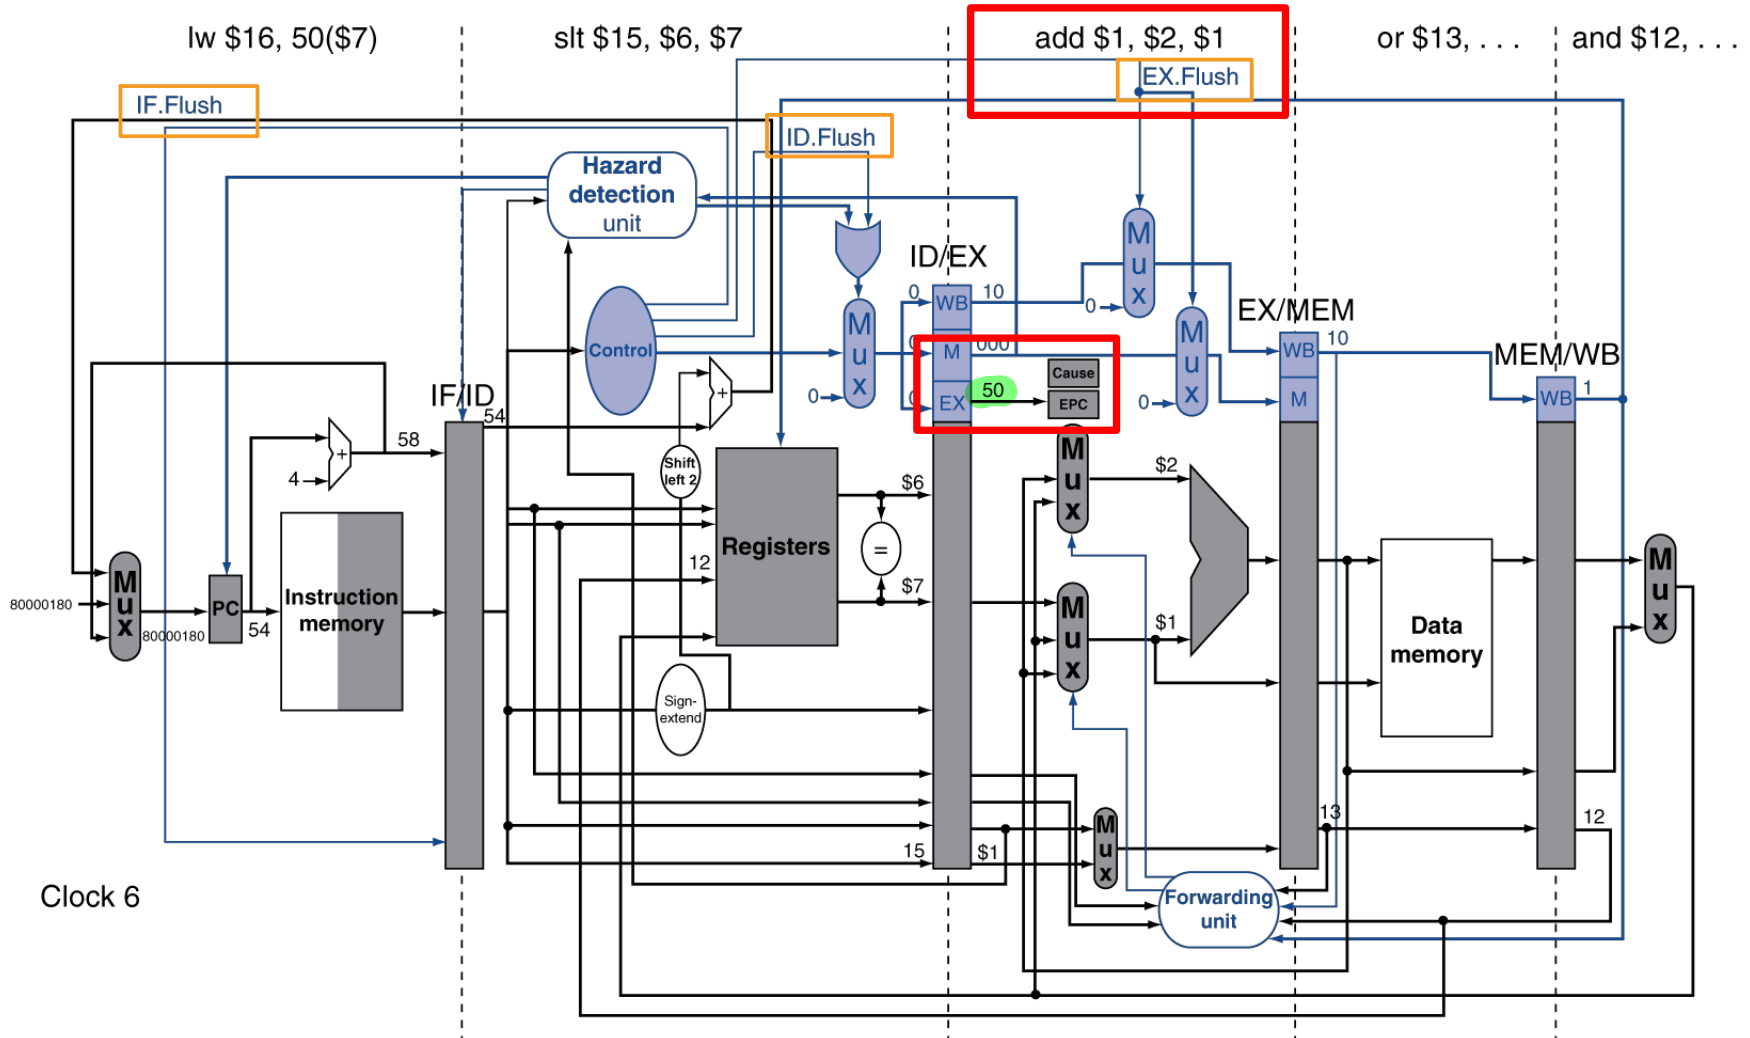
◆ Handler

80000180 sw \$25, 1000(\$0)

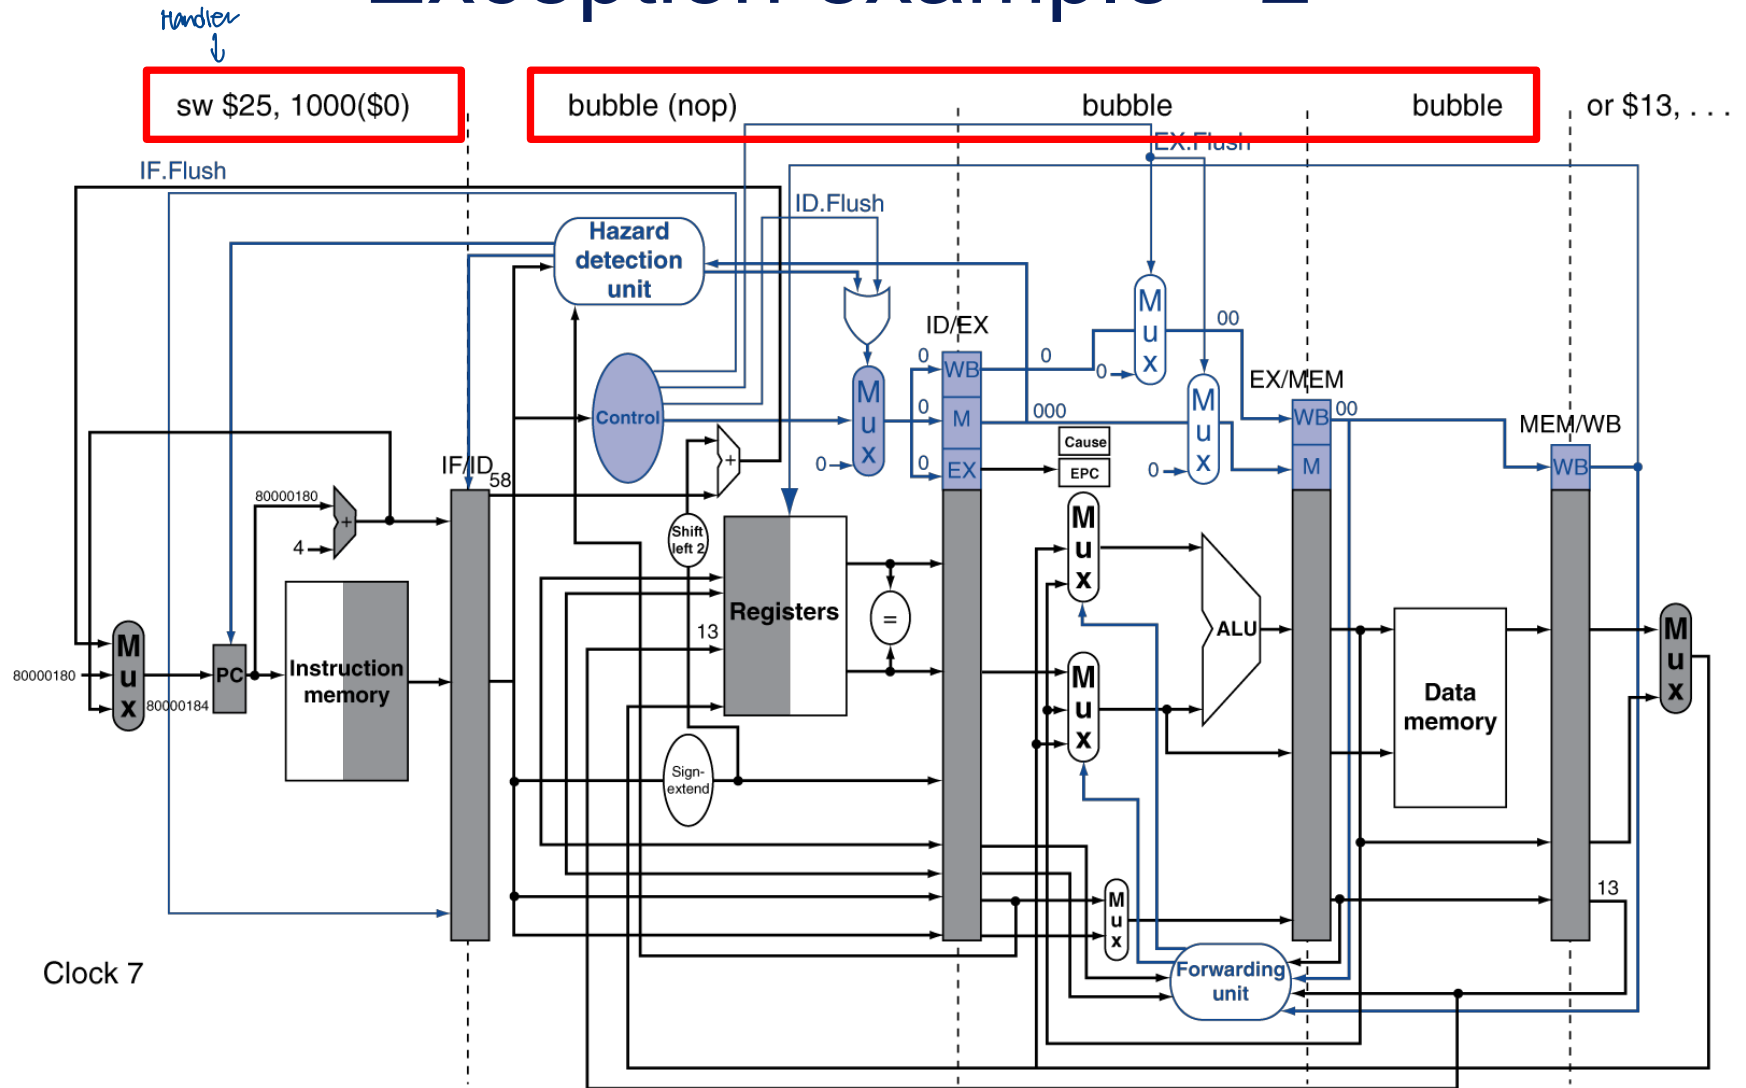
80000184 sw \$26, 1004(\$0)

...

Exception example - 1



Exception example - 2



Interrupt handler examples

- ◆ On asynchronous interrupts, **device-specific handlers** are invoked to service the I/O devices
- ◆ On exceptions, kernel handlers are invoked to either
 - Correct the faulting condition and continue the program (e.g., emulate the missing FP functionality), or
 - Signal back to the user process if a user-level handler function is registered, or
 - Kill the process if the exception cannot be corrected
- ◆ “System call” is a special kind of function call from user process to kernel-level service routines

Registering interrupt handlers

- ◆ You can register custom interrupt handlers

```
#include <stdio.h>
#include <signal.h>
#include <stdlib.h>

void my_handler(int sig) {
    printf("Caught signal %d: Handling FP exception!\n", sig);
    exit(EXIT_FAILURE); // Terminate the program
}

int main() {
    signal(SIGFPE, my_handler); // Register handler
    int x = 1 / 0; // This causes a divide-by-zero exception
    printf("This line will not execute.\n");
    return 0;
}
```


Returning from interrupts

- ◆ Software restores all architectural states saved at the start of the interrupt routine

(1) MIPS32 uses a special jump instruction

- ERET : MIPS 명령어

- > To atomically restore the automatically saved CPU states
- > To atomically restore the privilege level *kernel to user*
- > To atomically jump back to the interrupted address in EPC

(2) MIPS R2000 uses a pair of instructions

(assume “mfc0 r26,epc” done)

`jr r26` // jump to a copy of EPC in r26

`rfe` // restore from exception mode (control inst)

Another problem of delay slot

- ◆ What if the faulting address is in the delay slot?
 - Simply jumping back to EPC will cause error

```
beq $t0, $t1, target
```

```
lw $t2, $t3, 0($t4) // Branch Delay Slot (Cause exception)
```

...  Delay slot에 위치하는 명령어에서 Exception → EPC=EPC-4로 클리어되지 않음.

```
target:
```

```
add $t5, $t6, $t1 // This is what should be executed
```

- ◆ MIPS sets EPC = EPC - 4 to handle this type of errors

- Beq is executed twice, but does not affect the functionality

1) Would Exception → Exception Handler → kernel mode로 무조건 바뀔 수 있음.

→ Branch Condition이 달라질 수도 있다 → Branch부터 다시 실행하는 것이 안전함

Nesting Interrupts

- ◆ On an interrupt control transfer, further exceptions or asynchronous interrupts are **disabled automatically**
 - Another interrupt would overwrite the contents of the EPC and Interrupt Cause and Status Registers
 - (1) The handler must be carefully written not to generate synchronous exceptions itself during this window of vulnerability
 - (2) The handler must disable further asynchronous interrupts using interrupt status register (CR12)

→ Nesting: EPC, state registers가 덮어쓰일 수 있다.
→ Interrupt handler에서 exception, Interrupt가 발생하면 종료 ← MIPSE 막아야 함.
- ◆ However, for long-running handlers, interrupt must be enabled not to miss critical interrupts
 - Handler가 끝난 후에 Interrupt를 반드시 재-enable 해야 함
 - The handler must save the contents of EPC/Cause/Status to memory (stack) before re-enabling asynchronous interrupt
 - Once interrupts are re-enabled, EPC/Cause/Status can be updated by the next interrupt

Interrupt Priority

- ◆ Asynchronous interrupt sources are ordered by **priorities**
 - Higher-priorities interrupts are more timing critical
 - If multiple interrupts are triggered, the handler handles the highest-priority interrupt first
- ◆ Interrupts from different priorities can be selectively disabled by setting the mask in the Status register
- ◆ When servicing a particular priority interrupt, the handler only enable higher-priority interrupts
 - Higher-priority interrupt should not get delayed

Frequently enabling low-priority interrupts may critically affect the CPU performance (user-level instructions)

Short interrupt

```
_handler_shortest:
```

```
    // no prologue needed (Do not store registers)
```

```
    // use only r26 ($k0), r27 ($k1)
```

```
    ... Short handler body ...
```

```
    // Use only r26, r27
```

```
    // Blocks other interrupts
```

```
    // epilogue
```

```
    mfc0 $r26, $epc
```

```
    // Set r26 to epc
```

```
    jr $r26
```

```
    // restore PC
```

```
    rfe
```

```
    // restore from exception mode
```

Longer interrupt (Not nested)

_handler_longer:

// Prologue (save registers to use)

// use r8/9 (callee-saved), r26 (\$k0), r27 (\$k1)

addi \$sp, \$sp, -8 // Store r8/9

sw \$r8, 0(\$sp)

sw \$r9, 4(\$sp)

... longer handler body ... // Use only r8/9, r26, r27
// Blocks other interrupts

// epilogue

lw \$r8, 0(\$sp) // Restore r8/9

lw \$r9, 4(\$sp)

addi \$sp, \$sp, 8

mfc0 \$r26, \$epc // Set r26 to epc

jr \$r26 // restore PC

rfe // restore from exception mode

Nestable handler

```
_handler_nestable:
```

```
// Support nestable setup @ prologue
```

```
addi $sp, $sp, -4 // Store EPC (to the stack)
```

```
mfc0 $r26, $epc
```

```
sw $r26, 0($sp)
```

```
addi $r26, $r0, 0x405 // Enable nested interrupt
```

```
mtc0 $r26, $status
```

```
... longer handler body ... // These can be interrupted  
// Blocks other interrupts
```

```
// epilogue
```

```
addi $r26, $r0, 0x404 // disable nested interrupts
```

```
mtc0 $r26, $status
```

```
lw $r26, 0($sp) // Get EPC from the stack
```

```
addi $sp, $sp, 8
```

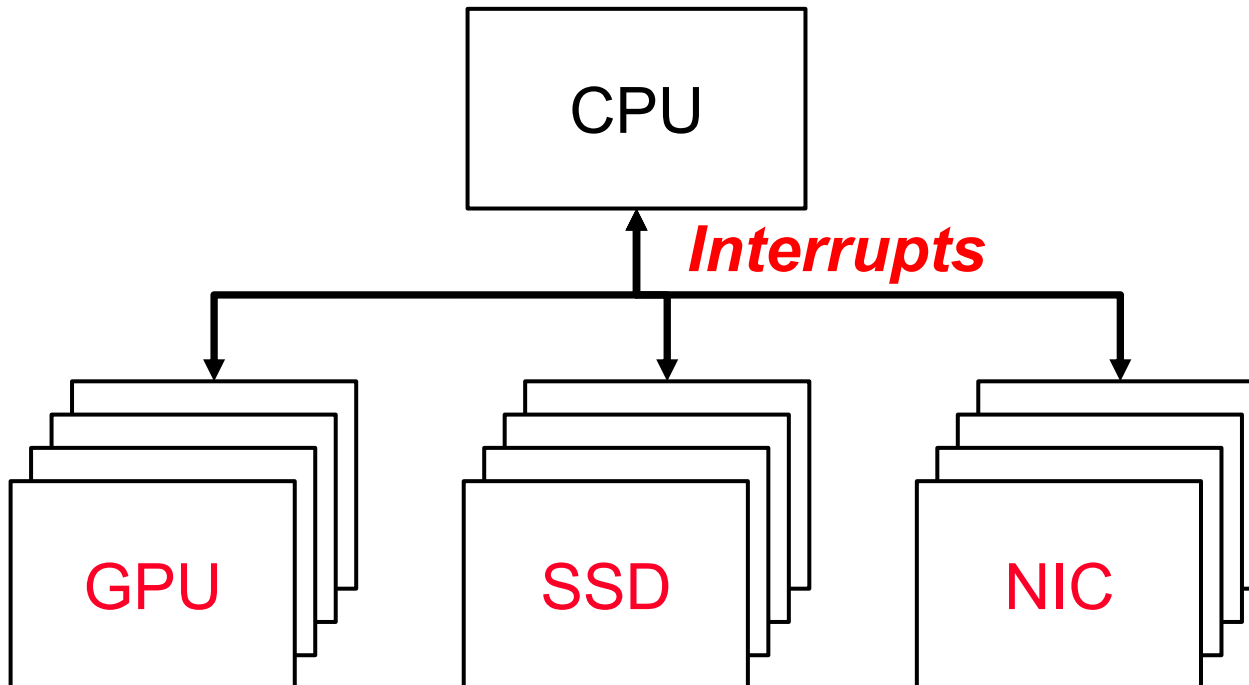
```
mfc0 r26, epc // Set r26 to epc
```

```
jr r26 // restore PC
```

```
rfe // restore from exception mode
```

Thinking of (external) interrupts

- ◆ The modern datacenter CPUs suffer from significantly large number of interrupts



Question?

Announcements:

Reading: finish reading P&H Ch.4

Handouts: none