

Lecture 20:

Multicore – Intro

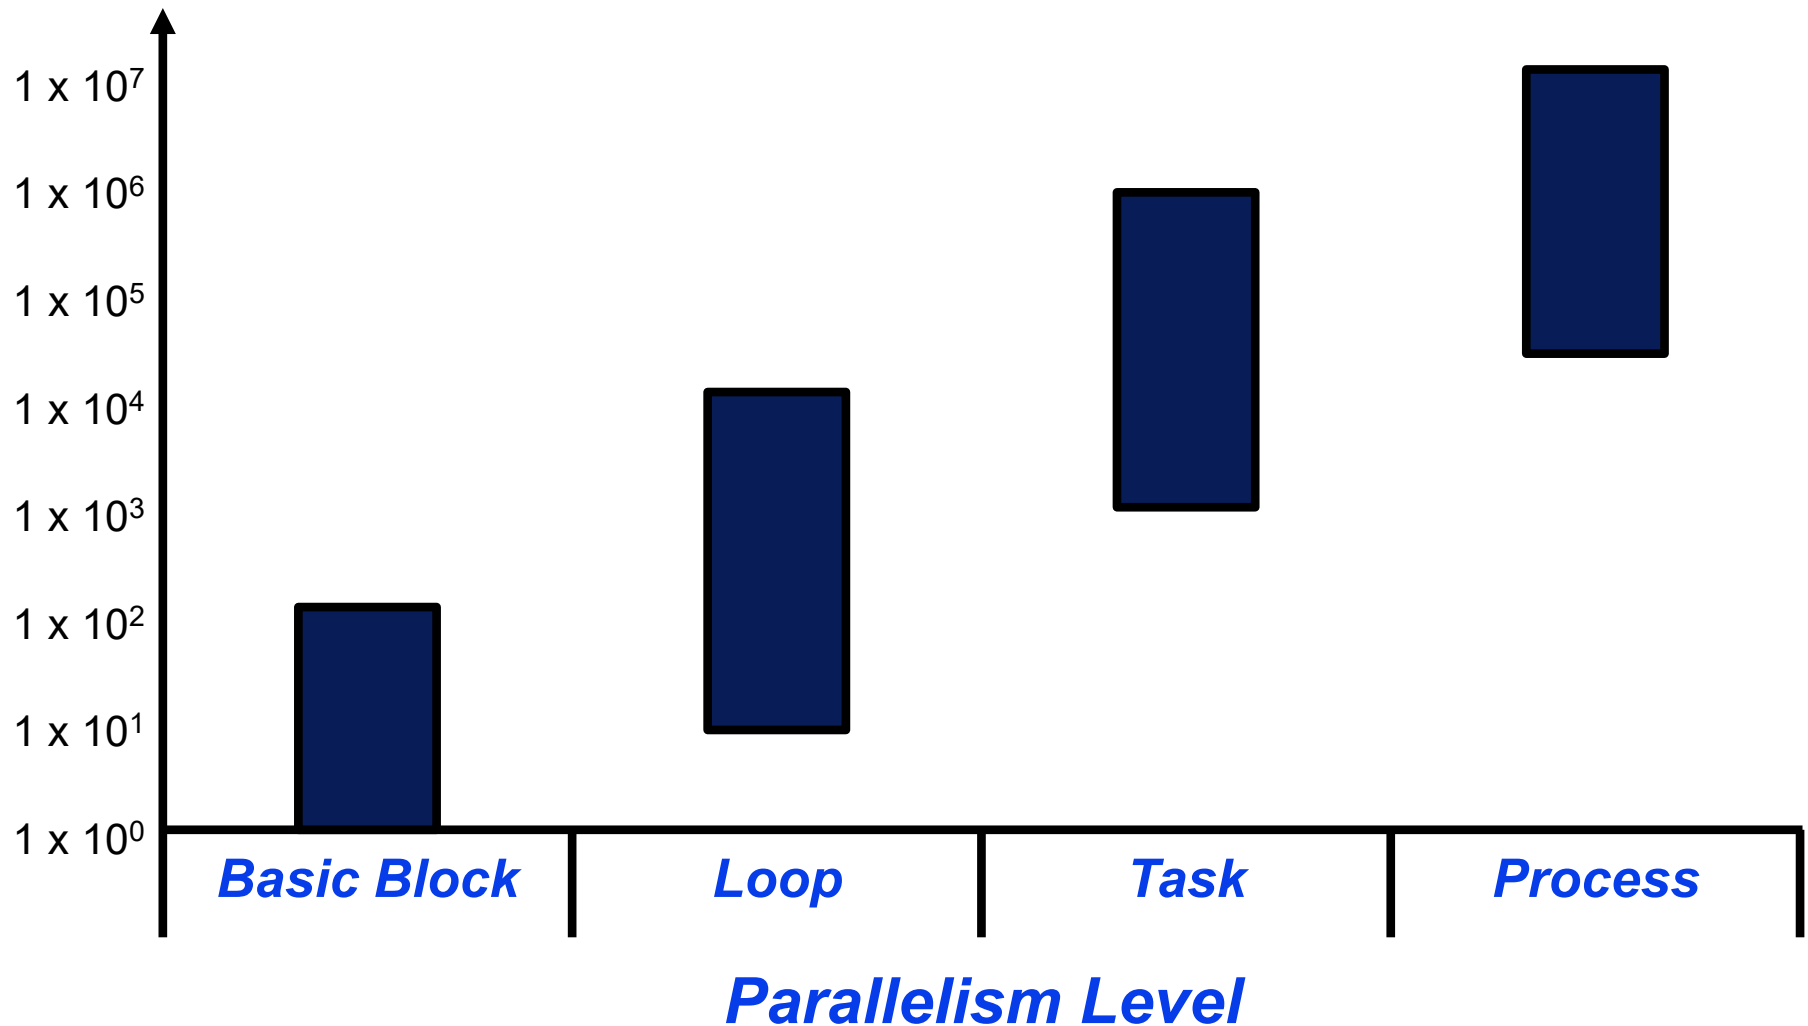
Hunjun Lee

hunjunlee@hanyang.ac.kr

Parallelism: the way to improve the performance!

- ◆ Instruction: all applications possess some parallelism
- ◆ Basic block: a group of instructions within a basic block (sequence of instructions between control flow operations)
- ◆ Loop iterations: each iteration of a loop may sometimes operate with independent data (run in parallel!)
- ◆ Tasks: large & independent functions in a single application (e.g., threads)
- ◆ Processes: independent OS processes


Parallelism: the way to improve the performance!




So far, ...

- ◆ You only considered that there is a single centralized processing unit
- ◆ This greatly simplifies the memory ... → The data is modified at a single centralized unit
- ◆ To exploit the parallelism, we need an extremely complex pipelining techniques to run multiple instructions in parallel
 - **Applies only to instructions, basic blocks, and a little bit of loop!**

Parallelism: the way to improve the performance!



Moore's Law




The performance scaling of a single core is reaching a dead end


One of the solutions ...

Symmetric Multiprocessor (SMP)

- ◆ We utilize multiple processors to exploit the parallelism! (Not a strict terminology though)
- ◆ Slow processor to processor communication (We cannot exploit the cache as the shared memory)




Parallelism: the way to improve the performance!

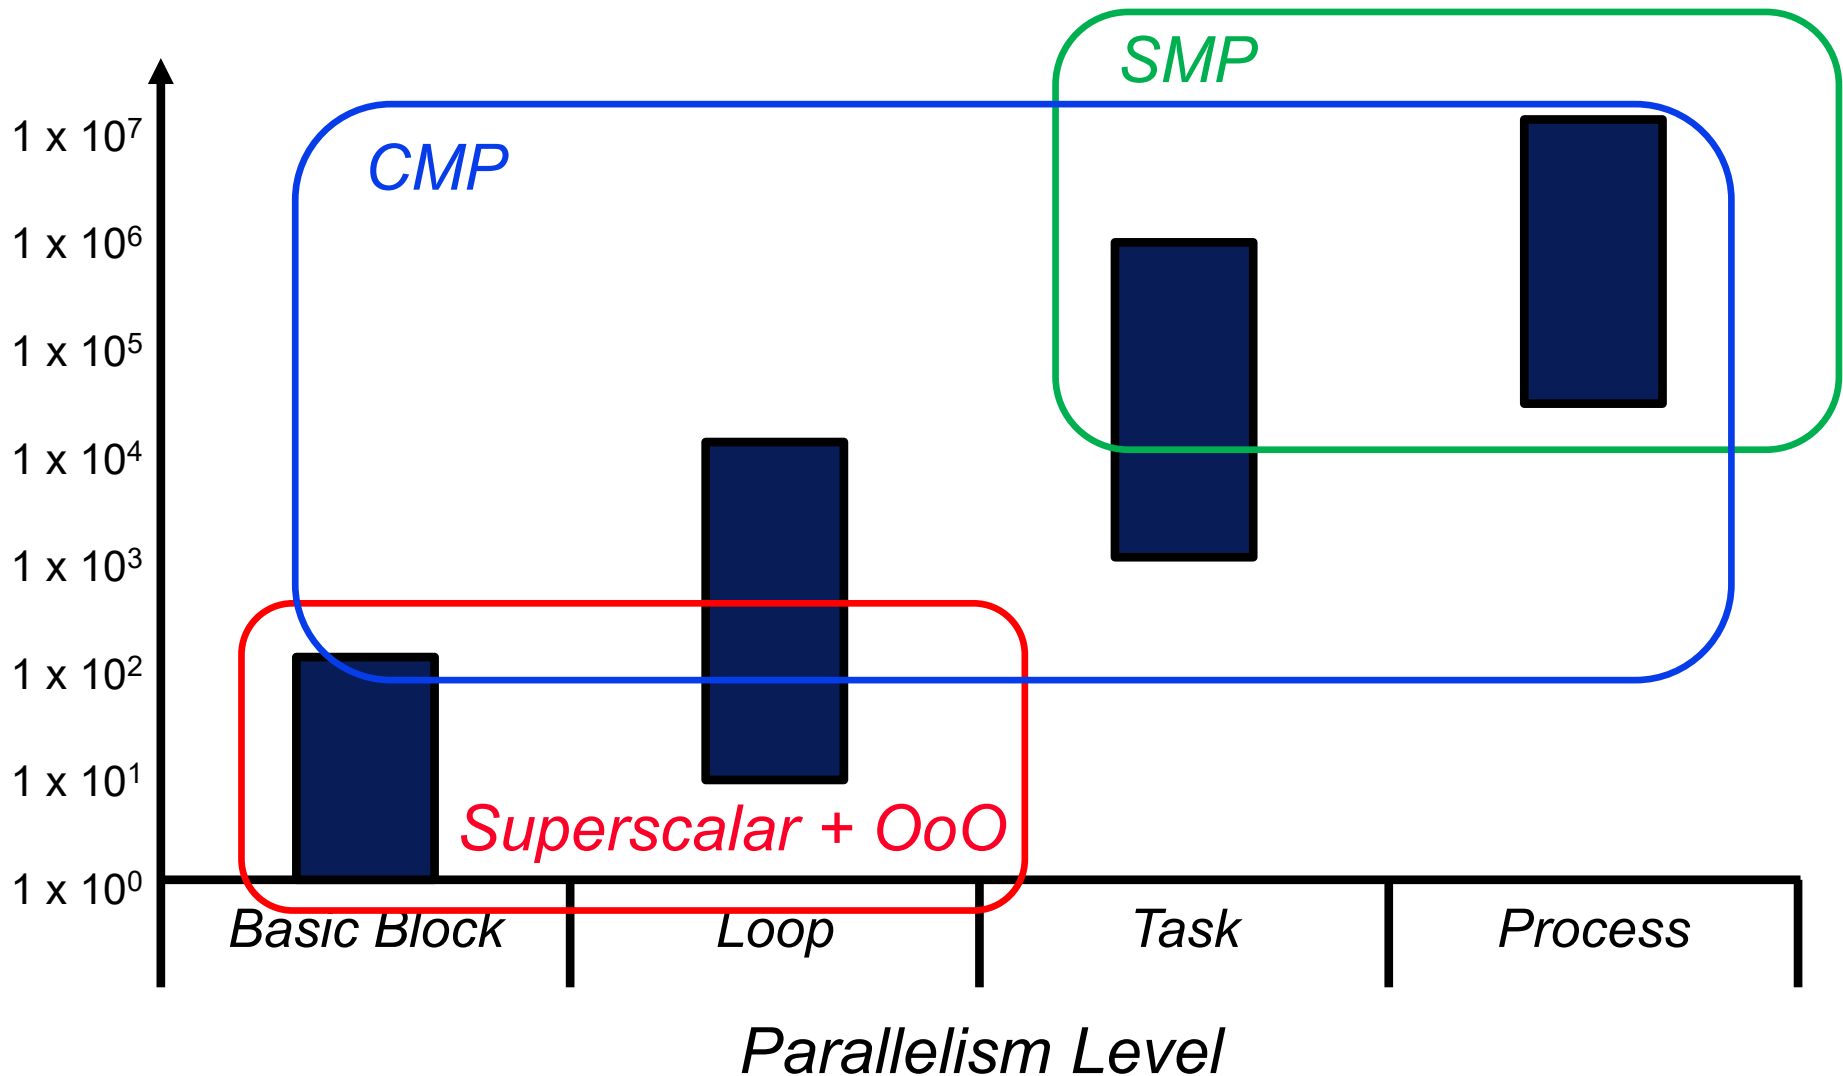


New Approach: The Chip Multiprocessor (CMP)

- ◆ Make a single processor which consists of multiple cores!
→ Allow faster data communication via a shared memory or an interconnection network
- ◆ A task is split across multiple cores → operating in parallel

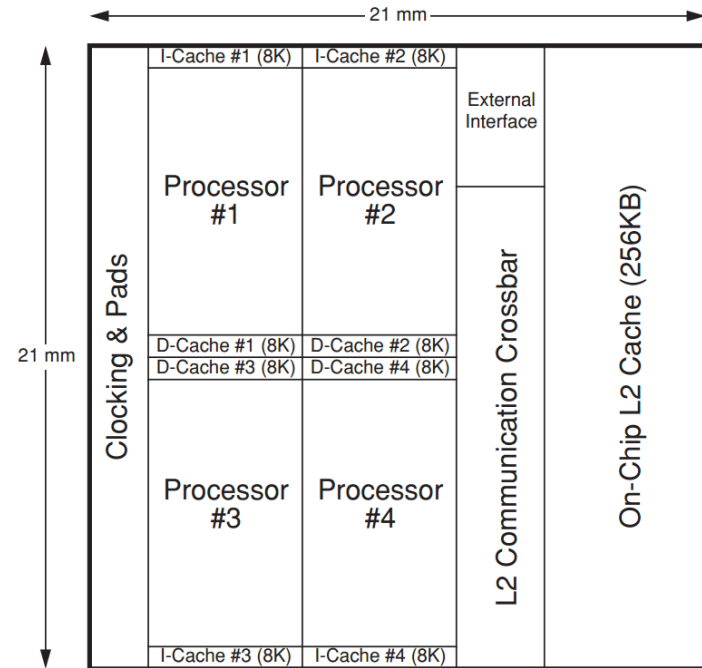
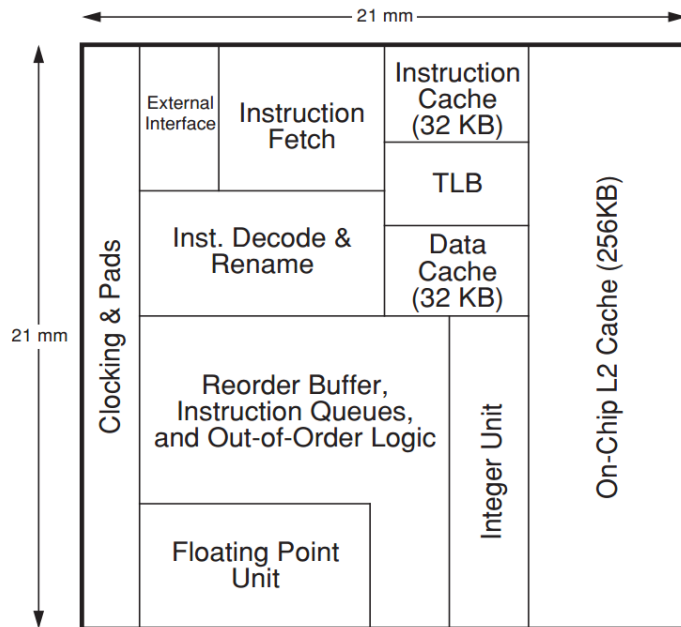


Parallelism: the way to improve the performance!



OoO vs. CMP

- ◆ 6-issue superscalar vs. 4-way CMP (1-way in-order core) → require similar amount of area ...
 - But 4-way CMP can exploit more distance parallelization opportunities



Partitioning Workloads

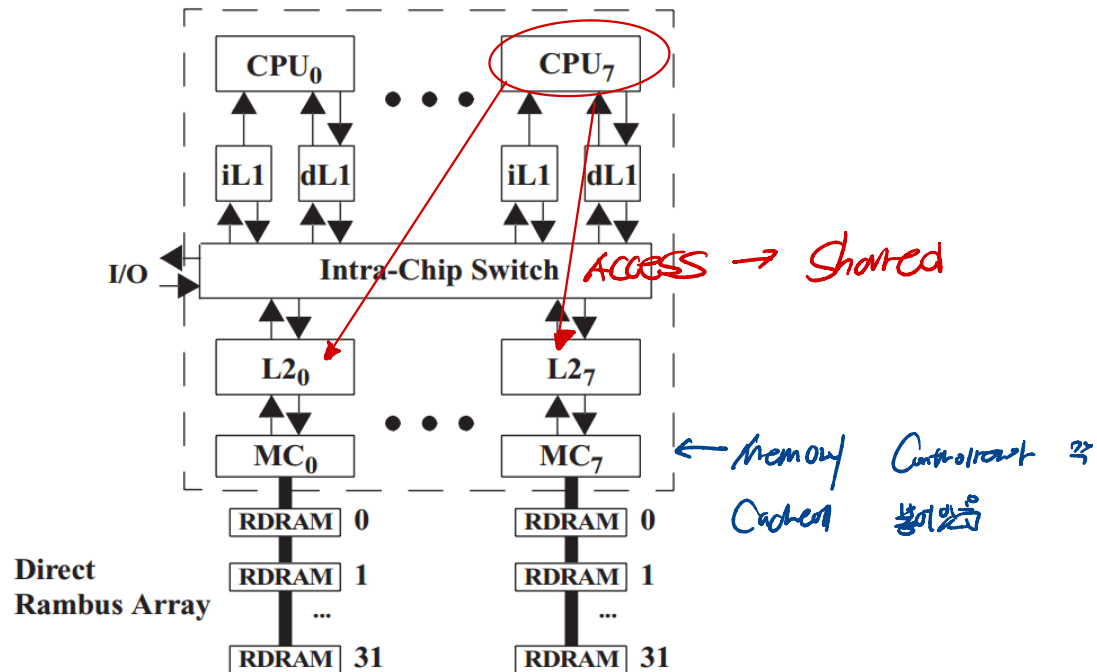
- ◆ We need to partition a single problem into multiple related tasks (threads)
- ◆ Loosely coupled multiprocessor
 - **No shared global memory address space**
 - Multicomputer network (network-based)
 - **Programmed via message passing (e.g., MPI) → explicit calls for communication**
- ◆ Tightly coupled multiprocessors
 - **Shared global memory address space**
 - Utilizing CMP
 - **Programming model similar to uniprocessors, but need to define operations on shared data and synchronization**

History of CMP – 1

Piranha Server CMP (ISCA'00)

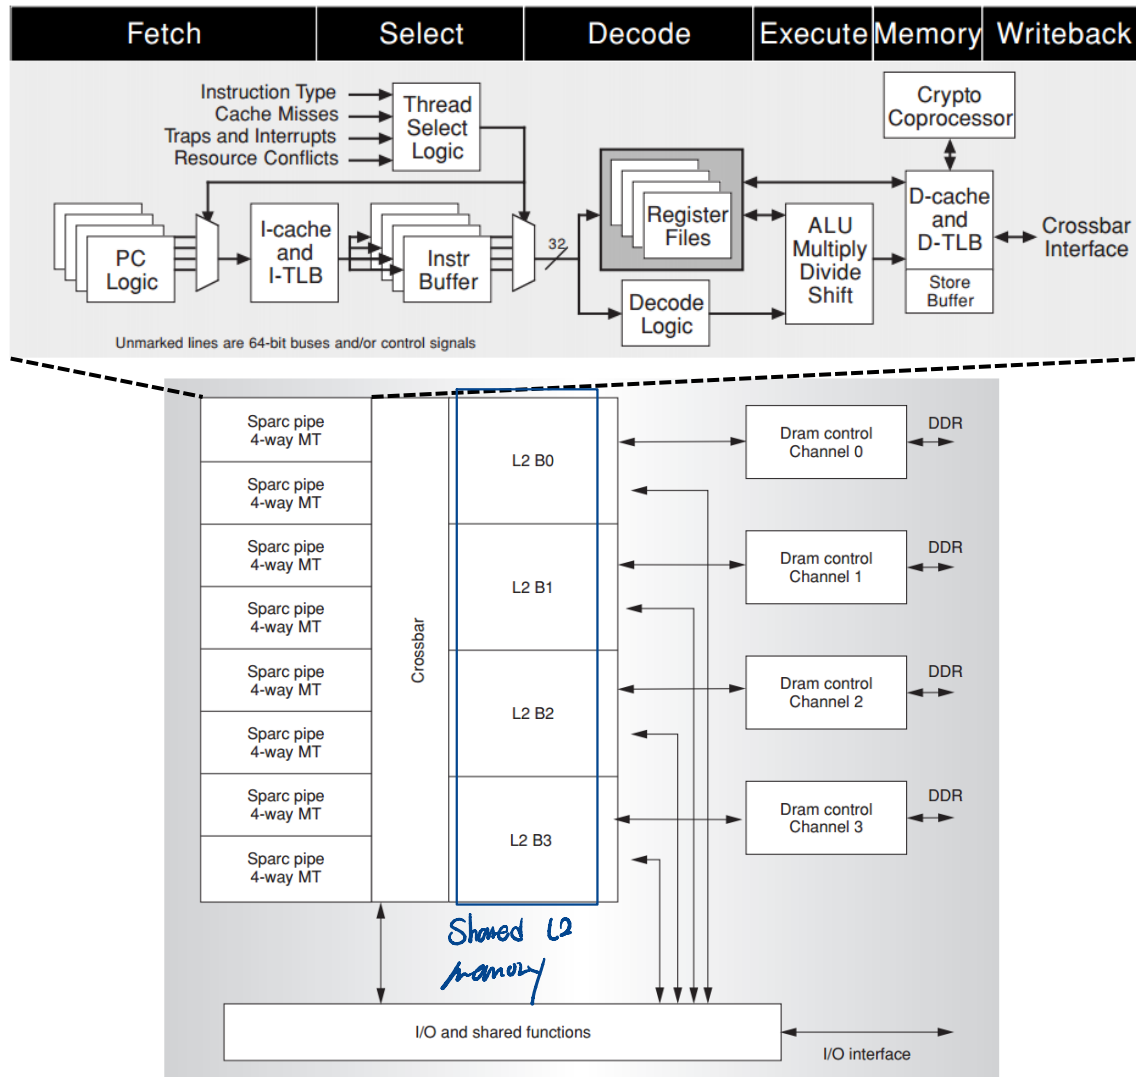
32:1 CMP Architecture

- ◆ Piranha was one of the first CMP architecture
 - Each core: 1-way in-order core * 8
 - Memory: Physically partitioned (but logically shared L2 cache)



History of CMP – 2

Niagara Processor (IEEE MICRO'05)



Performance implications of CMP

◆ Recall Amdahl's law

- f : parallelizable fraction of a program *CMP에서 병행적으로 실행할 수 있는 fraction*
- N : number of processors *프로세서 개수*
- Then, the ideal speedup would be ...
- $1 / (1 - f + f / N)$
- So, there's a clear limit on how fast CMP can become ...

◆ Also, you cannot achieve (f / N) even for parallelizable code

- Synchronization overhead (updates to the shared data)
- Load imbalance (the execution latency of each thread are not equal)
- Resource sharing overhead (such as memory bandwidth)

Multithreading

- ◆ Multiple processes/threads execute in parallel to speedup the overall task
- ◆ “Processes” get a unique virtual address space, and the entire states are copied
 - Including the memory, registers, etc ...
 - Forking: each process gets a unique VA which initially maps to the same PA (shared)
 - Apply copy-on-write (COW): Whenever there is a write, make a copy and change mapping!
- ◆ “Threads” share virtual address space
 - A “process” can comprise multiple “threads”
 - Threading: they share a common VA and map to a common PA, allocates a new (dedicated) stack space, new register context, ...

Question?

Announcements:

Reading: *read P&H Ch.6*

Handouts: *none*